

Yingjie Yang

LLM Post-Training | Agentic RL | Data Synthesis

📍 Beijing, China | 📞 +86 131-6121-0708 | ✉ thejayverse3@gmail.com | 🌐 jayverse3.github.io

Education

Peking University | Computer Science and Technology | M.S. Sept 2025 – June 2028
Northwest A&F University | Software Engineering | B.Eng. Sept 2020 – June 2024

Experience

UniPat AI | Research Intern Oct 2026 – present
 Work on high-quality training data synthesis and benchmark development for frontier coding agents.

Baidu | Basic Model R&D Department | Research Intern (**Top Internship Program**) Sept 2026 – Oct 2026
 Contributed to ERNIE's post-training, focusing on improving its performance in terminal environments and on Terminal-Bench.

ByteDance | AI Data & Safety | Research Intern Feb 2026 – Aug 2026
Collaborated with the Seed team to improve Seed models' ability to perform agentic coding and GUI tasks by sourcing, curating, and synthesizing training data from diverse sources. Analyzed training and evaluation results, rollout trajectories, and failure cases to identify model capability gaps, continuously refined data-synthesis pipelines, and **produced high-quality mid-training and post-training data at scale**.

• Coding Agent

- **Expert Data Curation and Generalization:** Contributed to annotation guideline development and guided domain experts in annotating tasks. Collected **2K+ expert-annotated tasks**, selected high-quality samples to build a seed pool, and generated **3K+ synthetic tasks** by generalizing from these seeds, bringing the candidate pool to **5K+ tasks**. Sampled five rollouts per task and selected verified trajectories to build RFT data, helping improve Seed models' score on **Terminal-Bench 2.0** from **45.2** to **53.6**.
- **Verifiable Task Synthesis from GitHub:** Built an automated task synthesis pipeline using Shell scripts scraped from GitHub. The pipeline generates instructions, environments, solutions, and tests in sequence, then iteratively repairs tasks based on feedback from Harbor Oracle / NOP validation and adversarial testing by a hack agent. To increase task difficulty and diversity, scraped Python source files, extracted complex workflows, and packaged them as Shell scripts, reusing the same pipeline to generate tasks. Sampled five rollouts per task and filtered out tasks with zero or five successful rollouts. Delivered **40K+ verifiable tasks** for Coding Agent post-training, helping improve Seed models' score on **Terminal-Bench 2.1** from **53.6** to **71.0**.
- **Targeted Data Synthesis for Failure Patterns:** Examined task requirements, reference solutions, and verification logic to understand the challenges in tasks where Seed models underperformed, then analyzed model trajectories to identify and categorize recurring failure patterns. Generated task proposals targeting model capability gaps, selected high-quality proposals through rubric-based scoring, and expanded them into complete tasks in the standard Harbor format. Conducted RFT using **2K+ purchased tasks** and **2K+ targeted synthetic tasks**, helping improve Seed models' score on **Terminal-Bench 2.1** from **71.0** to **84.7**.

• GUI Agent

- **GUI-Diver Dynamic Task Synthesis:** Selected 300 high-quality websites from [Similarweb](#) and [Tranco](#) and built GUI-Diver, an automated task synthesis pipeline. The pipeline first guides an agent to explore websites, select challenging scenarios, and generate task goals and action plans, then has the agent follow the plans to complete the tasks, collect trajectories, and generate full task descriptions and answers, and finally constructs corresponding verifiers.
- **Task Filtering and Training Validation:** Used model-as-judge to score synthesized tasks for quality and difficulty, selecting **300 high-quality tasks** to train **UI-TARS-1.5-7B** with GRPO and holding out 100 tasks for evaluation; the model's Pass@3 on this evaluation set improved from **38.7%** to **51.8%**. Delivered **900+ tasks** for GUI Agent post-training, helping improve Seed models' score on **OSWorld-Verified** from **61.9** to **78.8**.

Research

Seed2.1 Model Card: Agentic Intelligence for Productivity 2026
ByteDance Seed
Technical Report · [Project Page](#) · [Model Card](#)

Honors and Awards

- Kaggle Silver Medal: [LLMs - You Can't Please Them All](#) 2025
- Kaggle Silver Medal: [Learning Agency Lab - Automated Essay Scoring 2.0](#) 2024

Skills

Programming Languages: Python, C++ , TypeScript
Post-training: RLHF/RLVR, Process Reward Modeling, PPO, GRPO, DAPO, On-Policy Distillation
Framework: PyTorch, Transformers, LLaMA-Factory, verl, slime, Harbor
Tools: L^AT_EX, Git, Docker, Codex, Claude Code